

Cloud Computing Design Patterns

CIS437

Erik Fredericks // frederer@gvsu.edu

Adapted from Google Cloud Computing Foundations, Overview of Cloud Computing (Wufka & Canonico)

Design patterns for cloud apps!

<https://github.com/mehdihadeli/awesome-software-architecture/blob/main/docs/cloud-design-patterns/cloud-design-patterns.md>

<https://docs.aws.amazon.com/prescriptive-guidance/latest/cloud-design-patterns/introduction.html>


<https://www.techtarget.com/searchcloudcomputing/tip/5-cloud-design-patterns-to-create-resilient-applications>



First off...

What is a design pattern?

Next, some 350 slides to *remind us* what it is




What *is* a design pattern?

A design pattern “...names, abstracts, and identifies the key aspects of a common design structural that make it useful for creating a reusable object-oriented design.”*

A design pattern is a **proven** solution to a recurrent problem in a context.

An effective, reusable, proven structure/**communication** solution for a given object-oriented design problem.

What do we mean by proven?
How does communication fit in?



*From the book pictured

Why study them (the design patterns)

Reuse existing, high-quality solutions to commonly recurring problems

Establish common terminology to improve communications within teams

- Shifts the level of thinking to a higher perspective.

Improve team communication and individual learning

Improve modifiability and maintainability of code


- Design patterns are time-tested solutions (i.e., “proven”)


Why study them (the design patterns)

Adoption of improved object-oriented design strategies

- Encapsulation and information hiding
- Design to interfaces
- Favor composition over inheritance

```
class
{
    data members
    +
    methods (behavior)
}
```





Facade pattern

Pattern Category: **Structural**

Intent:

- Provide a unified interface to a set of interfaces in a subsystem.
- Facade defines a unified higher-level interface that makes the subsystems easier to use.

Problem addressed:

- Using design patterns often leads to a complex system of many small components which may be daunting for the casual user. It would be nice if there were a way to provide a simple interface for the basic functionality that is needed most often.

Solution:

- Create a Facade class that encapsulates the basic functionality of the system by bundling together common operations

When else would a *Facade* class be useful?

Cloud design patterns

Same as normal design patterns, but specific to cloud applications

- i.e., proven solutions to common problems

What are our concerns again?

- Normal application with:
 - Globally distributed userbase
 - Load spikes
 - ... etc.

Now...

There are a *lot* of design patterns out there

- And there are ever-growing lists for the cloud
 - <https://github.com/mehdihadeli/awesome-software-architecture/blob/main/docs/cloud-design-patterns/cloud-design-patterns.md>

We're going to walk through a few of them

- Keep learning though!
- A good portion of them can be useful **for your future career**

Cloud fallacies (similar to distributed computing fallacies)

- The network is reliable
- Latency is zero
- Bandwidth is infinite
- The network is secure
- Topology doesn't change
- There is one administrator
- Component versioning is simple
- Observability implementation can be delayed
 - i.e., monitoring and understanding what went wrong

https://en.wikipedia.org/wiki/Fallacies_of_distributed_computing

<https://learn.microsoft.com/en-us/azure/architecture/patterns/>

Publish/Subscribe

Asynchronous communication (decoupling transmission of data)

Publisher - sends out data on *topics*

Subscriber - receives data for *topics they've subscribed to*

(Broker) - *middleman to disseminate traffic / store who gets what*

Multiple choices of quality of service


- i.e., do we care if the data is received or non-corrupted?
 - Why?
 -

Common in IoT applications!

<https://docs.aws.amazon.com/prescriptive-guidance/latest/cloud-design-patterns/publish-subscribe.html>

<https://learn.microsoft.com/en-us/azure/architecture/patterns/publisher-subscriber>

Pub/sub




Clients and brokers

Client:

- Publisher or subscriber that connects to a **broker**
- Persistent (maintains connection) or transient (not tracked)

Broker (**central hub**):

- Receiving and filtering messages
- Understanding which clients are 'interested' in data
- Sending messages to subscribed clients
- Authenticating/authorizing clients



Topics

Hierarchical string that filters messages for clients

topic level
separator
↓
myhome / groundfloor / livingroom / temperature
└───┬───┘ └───┬───┘
topic level topic level

single-level
wildcard
↓
myhome / groundfloor / + / temperature
└───┬───┘ └───┬───┘
topic level topic level

only one level

multi-level
wildcard
↓
myhome / groundfloor / #
└───┬───┘ └───┬───┘
topic level topic level

only at the end

multiple topic levels

<https://www.youtube.com/watch?v=f5o4tlz2Zzc>

Geode



Geode

Why?

- Availability required worldwide (or, at least in multiple geographic regions)
- Scale required!

Concerns:

- Network latency + traffic management
- Worldwide deployment
- Data geo-distributed

Geode

How can we solve here?

Geode

How can we solve here?

- Create a bunch of geographic nodes (geodes...)
- "Satellite" deployments


Essentially, have a good devops approach (CI/CD)

- Deploy your app (ideally, templated)
- Reflect it to multiple regions automatically (via CI/CD)
- Load balance to direct traffic

Should never be used by itself (> 1 geode required for the pattern)

Updates to app reflect automatically to all geodes!

Geode




API Gateway

Target application:

- Collection of microservices
- Multiple client frontends

Similar to facade pattern!




API Gateway

Single point-of-entry from clients to backend


- Complexity behind the scenes hidden/abstracted

Difference to facade?

- Uses reverse proxy / gateway routing for communication
- i.e., requests from client are routed appropriately to microservice needed




How could we implement this?



How could we implement this?

Microservices with databases


- Naturally



Some form of routing application

- Combination of serverless functions that know where the microservices are?
- VM that handles reverse proxy (or a reverse proxy server itself)

Forward Proxy Flow



API Gateway

Advantages?

Disadvantages?

API Gateway

Advantages?

- Can aggregate client requests into single response
 - i.e., Multiple microservices queried and lumped together
- Load balancing possible
- Authorization/Authentication handled by networking layer

Disadvantages?

- Single(-ish) point of failure
- Extra complexity
- Possible anti-pattern
 - Bad design!
 - Could be giving the gateway "too much to do"

Circuit breaker

Prevents caller service from retrying after multiple timeouts/failures

- Detects when callee is available again

Possible causes:

- Network disruption
- Callee overloaded
- etc.

Issue avoiding?

- Consuming resources from numerous retried calls
 - Could impact cost and/or performance!


Circuit breaker



Circuit breaker with payment service failure



Circuit breaker stops routing to payment service



Circuit breaker periodically retries payment service

Circuit breaker

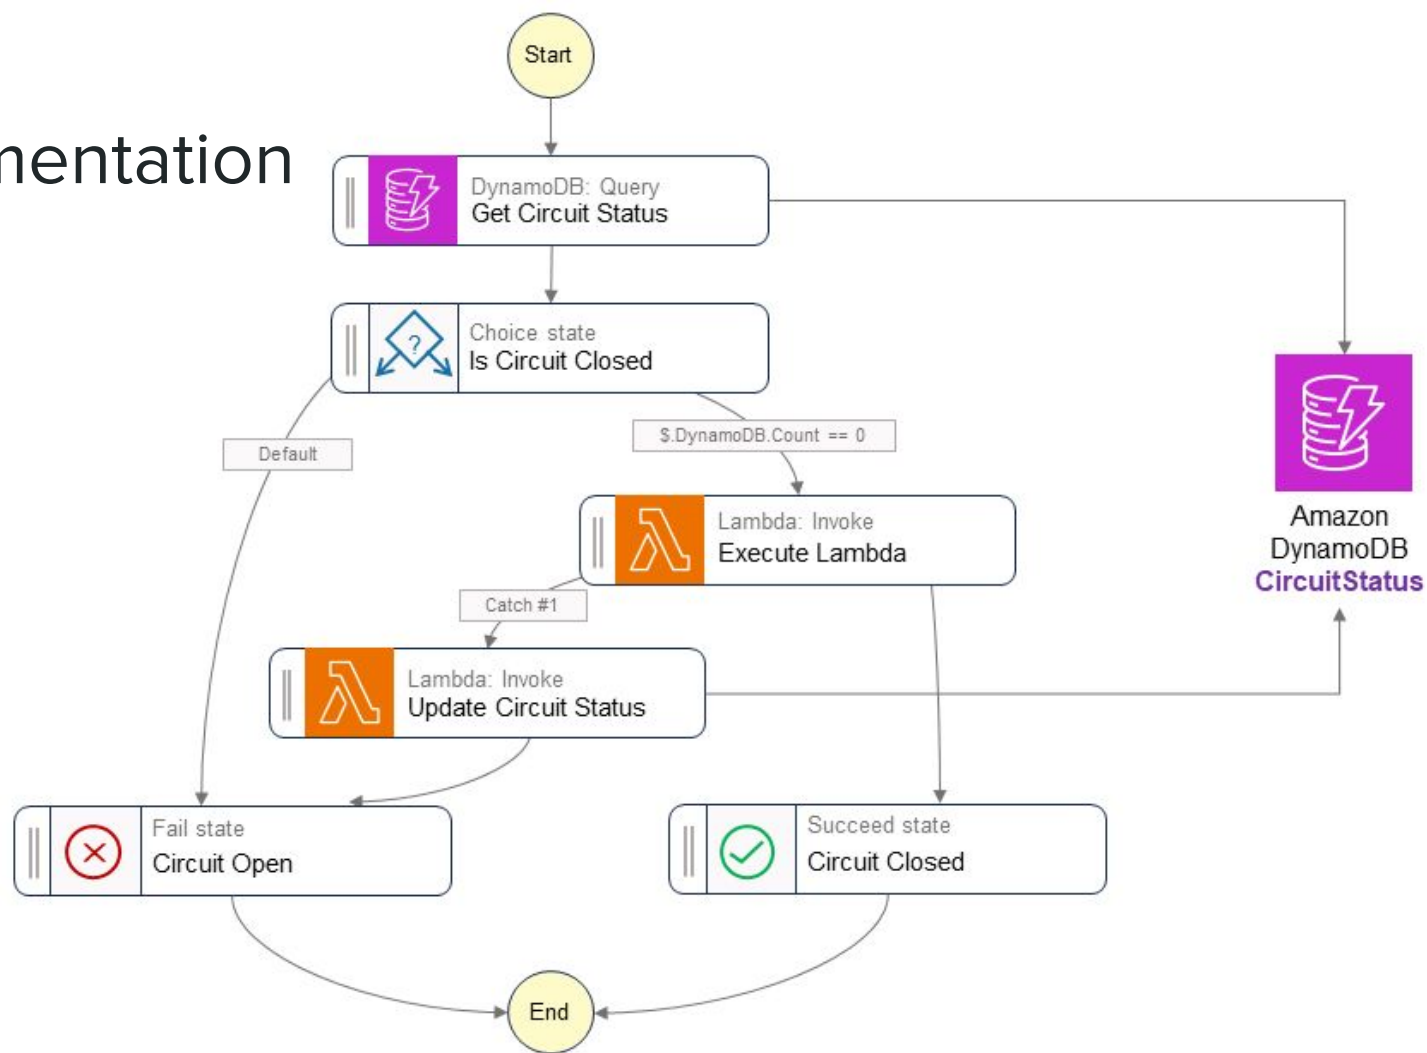


Circuit breaker periodically retries payment service



Circuit breaker with working payment service

AWS Implementation



And?

Advantages?

Disadvantages?

And?

Advantages?

- Reduction in unnecessary retry calls
- Possible reduction in 'stale' or duplicate calls
 - Perhaps a credit card auth. got stuck in the system?

Disadvantages?

- Complexity!
 - Requires multiple services (database, serverless, state machine, etc.)
 - Extra \$\$ for extra services!
- Multiple points of failure
 - What if you have an issue with your database now? Or one of your lambdas?